

Final Project Submission

Please fill out:

- Student name: Jane Martha
- Student pace: Part Time
- Scheduled project review date/time: 2/18/2024
- Instructor name: Samwel G
- Blog post URL:

```
# Your code here - remember to use markdown cells for comments as well!
```

```
# Import standard packages
```

```
import pandas as pd
import numpy as np
```

```
# import
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
%matplotlib inline
```

```
#loading data
```

```
df_title_ratings =
pd.read_csv("zippedData/imdb.title.ratings.csv.gz",)
df_title_basics = pd.read_csv("zippedData/imdb.title.basics.csv.gz",)
df_movie_gross = pd.read_csv("zippedData/bom.movie_gross.csv.gz",)
df_movies= pd.read_csv("zippedData/tmdb.movies.csv.gz",index_col=0)
```

```
#accessing the first 5 rows in df_title_basics
```

```
df_title_basics.head()
```

	tconst	primary_title
original_title \		
0	tt0063540	Sunghursh
1	tt0066787	One Day Before the Rainy Season
2	tt0069049	The Other Side of the Wind
3	tt0069204	Sabse Bada Sukh
4	tt0100275	The Wandering Soap Opera

	start_year	runtime_minutes	genres
0	2013	175.0	Action, Crime, Drama
1	2019	114.0	Biography, Drama
2	2018	122.0	Drama

3	2018	NaN	Comedy,Drama
4	2017	80.0	Comedy,Drama,Fantasy

#Checking metadata of df_title_basics

df_title_basics.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 146144 entries, 0 to 146143
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tconst                 146144 non-null object
1   primary_title          146144 non-null object
2   original_title         146123 non-null object
3   start_year             146144 non-null int64
4   runtime_minutes        114405 non-null float64
5   genres                 140736 non-null object
dtypes: float64(1), int64(1), object(4)
memory usage: 6.7+ MB
```

#accessing the last 5 rows in df_title_basics

df_title_basics.tail()

	tconst	primary_title \
146139	tt9916538	Kuambil Lagi Hatiku
146140	tt9916622	Rodolpho Teóphilo - O Legado de um Pioneiro
146141	tt9916706	Dankyavar Danka
146142	tt9916730	6 Gunn
146143	tt9916754	Chico Albuquerque - Revelações

	original_title	start_year \
146139	Kuambil Lagi Hatiku	2019
146140	Rodolpho Teóphilo - O Legado de um Pioneiro	2015
146141	Dankyavar Danka	2013
146142	6 Gunn	2017
146143	Chico Albuquerque - Revelações	2013

	runtime_minutes	genres
146139	123.0	Drama
146140	NaN	Documentary
146141	NaN	Comedy
146142	116.0	NaN
146143	NaN	Documentary

#Working with title_basics_df

df_title_basics.describe()

	start_year	runtime_minutes
count	146144.000000	114405.000000
mean	2014.621798	86.187247
std	2.733583	166.360590

min	2010.000000	1.000000
25%	2012.000000	70.000000
50%	2015.000000	87.000000
75%	2017.000000	99.000000
max	2115.000000	51420.000000

#Checking out the rows and columns of df_title_basics

df_title_basics.shape

(146144, 6)

checking the data types

df_title_basics.dtypes

```
tconst      object
primary_title  object
original_title object
start_year   int64
runtime_minutes float64
genres       object
dtype: object
```

#Accessing columns of df_title_basics

df_title_basics.columns

```
Index(['tconst', 'primary_title', 'original_title', 'start_year',
       'runtime_minutes', 'genres'],
      dtype='object')
```

#Accessing the index of df_title_basics

df_title_basics.index

RangeIndex(start=0, stop=146144, step=1)

#aceeing the first 5 rows in df_title_ratings

df_title_ratings.head()

	tconst	averagerating	numvotes
0	tt10356526	8.3	31
1	tt10384606	8.9	559
2	tt1042974	6.4	20
3	tt1043726	4.2	50352
4	tt1060240	6.5	21

#Accessing concise summary of df_title_ratings

df_title_ratings.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 73856 entries, 0 to 73855
Data columns (total 3 columns):
```

#	Column	Non-Null Count	Dtype
0	tconst	73856 non-null	object
1	averagerating	73856 non-null	float64
2	numvotes	73856 non-null	int64

dtypes: float64(1), int64(1), object(1)
memory usage: 1.7+ MB

#Accessing the last 3 rows of df_title_ratings.tail
df_title_ratings.tail(3)

	tconst	averagerating	numvotes
73853	tt9851050	4.7	14
73854	tt9886934	7.0	5
73855	tt9894098	6.3	128

#accessing description of df_movie_gross
df_title_ratings.describe()

	averagerating	numvotes
count	73856.000000	7.385600e+04
mean	6.332729	3.523662e+03
std	1.474978	3.029402e+04
min	1.000000	5.000000e+00
25%	5.500000	1.400000e+01
50%	6.500000	4.900000e+01
75%	7.400000	2.820000e+02
max	10.000000	1.841066e+06

#Analyzing rows and columns of df_title_ratings
df_title_ratings.shape

(73856, 3)

#Checking the data type
df_title_ratings.dtypes

tconst	object
averagerating	float64
numvotes	int64

dtype: object

#Accessing columns of df_title_ratings
df_title_ratings.columns

Index(['tconst', 'averagerating', 'numvotes'], dtype='object')

#Accessing the index of df_title_ratings
df_title_ratings.index

RangeIndex(start=0, stop=73856, step=1)

```
#Analyzing metadata of df_movie_gross.info
```

```
df_movie_gross.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 3387 entries, 0 to 3386  
Data columns (total 5 columns):  
#   Column                Non-Null Count  Dtype    
---  ---                  
0    title                 3387 non-null   object   
1    studio                3382 non-null   object   
2    domestic_gross        3359 non-null   float64  
3    foreign_gross         2037 non-null   object   
4    year                  3387 non-null   int64    
dtypes: float64(1), int64(1), object(3)  
memory usage: 132.4+ KB
```

```
#checking the first 5 rows of df_movie_gross
```

```
df_movie_gross.head()
```

	title	studio	domestic_gross
0	Toy Story 3	BV	415000000.0
1	Alice in Wonderland (2010)	BV	334200000.0
2	Harry Potter and the Deathly Hallows Part 1	WB	296000000.0
3	Inception	WB	292600000.0
4	Shrek Forever After	P/DW	238700000.0

	foreign_gross	year
0	652000000	2010
1	691300000	2010
2	664300000	2010
3	535700000	2010
4	513900000	2010

```
#checking the last 5 rows of df_movie_gross
```

```
df_movie_gross.tail()
```

	title	studio	domestic_gross
foreign_gross \			
3382	The Quake	Magn.	6200.0
NaN			
3383	Edward II (2018 re-release)	FM	4800.0
NaN			
3384	El Pacto	Sony	2500.0
NaN			

3385		The Swan	Synergetic	2400.0
NaN				
3386		An Actor Prepares	Grav.	1700.0
NaN				

	year
3382	2018
3383	2018
3384	2018
3385	2018
3386	2018

```
#print numbers pf rows and columns in bom_movie_gross_df
print(df_movie_gross.shape)
```

```
(3387, 5)
```

```
#Accessing the columns of df_movie_gross
print(df_movie_gross.columns)
```

```
Index(['title', 'studio', 'domestic_gross', 'foreign_gross', 'year'],
      dtype='object')
```

```
print(df_movie_gross.dtypes)
```

title	object
studio	object
domestic_gross	float64
foreign_gross	object
year	int64

```
dtype: object
```

```
#Accessing the index of df_movie_gross
print(df_movie_gross.index)
```

```
RangeIndex(start=0, stop=3387, step=1)
```

```
#Removing leading and trailing whitespace in df_title_basics columns
[col.strip() for col in df_title_basics.columns]
```

```
['tconst',
 'primary_title',
 'original_title',
 'start_year',
 'runtime_minutes',
 'genres']
```

```
#Checking out the missing values in df_title_basics
missing_values = df_title_basics.isna().sum()
missing_values
```

```
tconst          0
primary_title   0
original_title  21
start_year      0
runtime_minutes 31739
genres          5408
dtype: int64
```

#Checking out the missing values in df_title_basics

```
missing_values = df_title_ratings.isna().sum()
missing_values
```

```
tconst          0
averagerating    0
numvotes         0
dtype: int64
```

#Removing leading and trailing whitespace in

```
[col.strip() for col in df_movie_gross.columns]
```

```
['title', 'studio', 'domestic_gross', 'foreign_gross', 'year']
```

#Checking out the missing values in df_bom_movie_gross

```
missing_values = df_movie_gross.isna().sum()
missing_values
```

```
title          0
studio         5
domestic_gross 28
foreign_gross  1350
year           0
dtype: int64
```

```
df_movies.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 26517 entries, 0 to 26516
```

```
Data columns (total 10 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	26517 non-null	int64
1	genre_ids	26517 non-null	object
2	id	26517 non-null	int64
3	original_language	26517 non-null	object
4	original_title	26517 non-null	object
5	popularity	26517 non-null	float64
6	release_date	26517 non-null	object
7	title	26517 non-null	object
8	vote_average	26517 non-null	float64
9	vote_count	26517 non-null	int64

```
dtypes: float64(2), int64(3), object(5)
```

```
memory usage: 2.0+ MB
```

```
df_movies.head()
```

	genre_ids	id	original_language	\
0	[12, 14, 10751]	12444	en	
1	[14, 12, 16, 10751]	10191	en	
2	[12, 28, 878]	10138	en	
3	[16, 35, 10751]	862	en	
4	[28, 878, 12]	27205	en	

	original_title	popularity	release_date	\
0	Harry Potter and the Deathly Hallows: Part 1	33.533	2010-11-19	
1	How to Train Your Dragon	28.734	2010-03-26	
2	Iron Man 2	28.515	2010-05-07	
3	Toy Story	28.005	1995-11-22	
4	Inception	27.920	2010-07-16	

	title	vote_average	vote_count
0	Harry Potter and the Deathly Hallows: Part 1	7.7	10788
1	How to Train Your Dragon	7.7	7610
2	Iron Man 2	6.8	12368
3	Toy Story	7.9	10174
4	Inception	8.3	22186

```
df_movies.tail()
```

	Unnamed: 0	genre_ids	id	original_language	\
26512	26512	[27, 18]	488143	en	
26513	26513	[18, 53]	485975	en	
26514	26514	[14, 28, 12]	381231	en	
26515	26515	[10751, 12, 28]	366854	en	
26516	26516	[53, 27]	309885	en	

	original_title	popularity	release_date	title	\
26512	Laboratory Conditions	0.6	2018-10-13	Laboratory	


```

Conditions
26513      _EXHIBIT_84xxx_      0.6    2018-05-01
_EXHIBIT_84xxx_
26514      The Last One      0.6    2018-10-01      The
Last One
26515      Trailer Made      0.6    2018-06-22
Trailer Made
26516      The Church      0.6    2018-10-05      The
Church

```

```

      vote_average  vote_count
26512           0.0           1
26513           0.0           1
26514           0.0           1
26515           0.0           1
26516           0.0           1

```

```
df_movies.shape
```

```
(26517, 10)
```

```
df_movies.describe()
```

```

      Unnamed: 0      id  popularity  vote_average
vote_count
count  26517.000000  26517.000000  26517.000000  26517.000000
26517.000000
mean   13258.000000  295050.153260    3.130912    5.991281
194.224837
std     7654.94288   153661.615648    4.355229    1.852946
960.961095
min      0.000000    27.000000    0.600000    0.000000
1.000000
25%     6629.000000  157851.000000    0.600000    5.000000
2.000000
50%     13258.000000  309581.000000    1.374000    6.000000
5.000000
75%     19887.000000  419542.000000    3.694000    7.000000
28.000000
max     26516.000000  608444.000000   80.773000   10.000000
22186.000000

```

```
df_movies.dtypes
```

```

Unnamed: 0      int64
genre_ids      object
id             int64
original_language  object
original_title  object
popularity     float64
release_date   object

```

```

title          object
vote_average   float64
vote_count     int64
dtype: object

df_movies.columns

Index(['Unnamed: 0', 'genre_ids', 'id', 'original_language',
      'original_title',
      'popularity', 'release_date', 'title', 'vote_average',
      'vote_count'],
      dtype='object')

df_movies.index

RangeIndex(start=0, stop=26517, step=1)

#Removing leading and trailing whitespace in df_title_ratings columns
[col.strip() for col in df_title_ratings.columns]

['tconst', 'averagerating', 'numvotes']

#Calculating the percentage of missing values per column with respect
to the entering df a function
def missing_values(data):
    """A simple function to identify data with missing values"""
    #identify the total missing values per column
    #sort in order
    miss = data.isnull().sum().sort_values(ascending = False)

    #calculate percentage of missing values
    percentage_miss = (data.isnull().sum() /
len(data)).sort_values(ascending = False)

    #store in a dataframe
    missing = pd.DataFrame({"Missing Values": miss, "Percentage":
percentage_miss}).reset_index()

    return missing
#applying function to the df_title_basics
missing_data = missing_values(df_title_basics)
missing_data

```

	index	Missing Values	Percentage
0	runtime_minutes	31739	0.217176
1	genres	5408	0.037005
2	original_title	21	0.000144
3	start_year	0	0.000000
4	primary_title	0	0.000000
5	tconst	0	0.000000

#genres and original_title missing values are few hence drop them without causing any effect on data

#Filtering df that only contains Nan and empty strings

```
missing_genres = df_title_basics.genres.isna()
df_title_basics = df_title_basics[~missing_genres]
```

```
missing_title = df_title_basics.original_title.isna()
df_title_basics = df_title_basics[~missing_title]
```

#Confirming whether Nan values in original_title and genres columns have all dropped

```
missing_values(df_title_basics)
```

	index	Missing Values	Percentage
0	runtime_minutes	28502	0.202524
1	genres	0	0.000000
2	start_year	0	0.000000
3	original_title	0	0.000000
4	primary_title	0	0.000000
5	tconst	0	0.000000

#Confirming whether Nan values in genres columns have all dropped

```
missing_values(df_title_basics)
```

	index	Missing Values	Percentage
0	runtime_minutes	28503	0.202528
1	original_title	2	0.000014
2	genres	0	0.000000
3	start_year	0	0.000000
4	primary_title	0	0.000000
5	tconst	0	0.000000

#Visualization of runtime minutes

```
runtimes = df_title_basics.runtime_minutes
```

#finding max and min runtime

```
min_runtime = runtimes.min()
max_runtime = runtimes.max()
mean_runtime = runtimes.mean()
print(f"minimum runtime: {min_runtime}")
print(f"Maximum runtime: {max_runtime}")
print(f"Mean runtime: {mean_runtime}")
```

#Choosing boxplot column


```
col_data = df_title_basics.runtime_minutes
```

Creating boxplot

```
plt.figure(figsize=(12,6))
```

```
sns.set_context('notebook')
sns.boxplot(x= col_data, color= "yellow")
plt.title('Runtime minutes Boxplot');
```

```
minimum runtime: 1.0
Maximum runtime: 51420.0
Mean runtime: 86.26155641884668
```



#Checking for the highest runtime


```
df_title_basics.loc[df_title_basics.runtime_minutes == max_runtime]
```

	tconst	primary_title	original_title	start_year
runtime_minutes \				
132389	tt8273150	Logistics	Logistics	2012
51420.0				

	genres
132389	Documentary

```
sns.pairplot(df_movie_gross)
```

```
<seaborn.axisgrid.PairGrid at 0x2a005f73fd0>
```



```
#replacing the missing values with median runtime
df_title_basics.runtime_minutes.fillna(df_title_basics.runtime_minutes
                                          .median(), inplace=True)
```

```
#checking for residual of missing values in DF
missing_values
```

```
title          0
studio         5
domestic_gross 28
foreign_gross  1350
year           0
dtype: int64
```

```
#checking duplicates
```

```
duplicate_rows = df_title_basics.duplicated().sum()
print(f"Num of Duplicated Rows: {duplicate_rows}")
```

```
Num of Duplicated Rows: 0
```

```
#Dealing with df_movie_gross
```

```
missing_data = missing_values(df_movie_gross)
missing_data
```

	index	Missing Values	Percentage
0	foreign_gross	1350	0.398583
1	domestic_gross	28	0.008267
2	studio	5	0.001476
3	year	0	0.000000
4	title	0	0.000000

```
#Filtering dataframe to remove missing values in df_movie_gross
```

```
missing_studio = df_movie_gross.studio.isna()
df_movie_gross = df_movie_gross[~missing_studio]
```

```
#checking the residual of missing values
```

```
missing_values(df_movie_gross)
```

	index	Missing Values	Percentage
0	foreign_gross	1349	0.398876
1	domestic_gross	26	0.007688
2	year	0	0.000000
3	studio	0	0.000000
4	title	0	0.000000

```
#removing unwanted characters like commas etc
```

```
df_movie_gross.foreign_gross.replace(',', '', inplace=True, regex=True)
```

```
df_movie_gross.foreign_gross =
df_movie_gross.foreign_gross.astype('float64')
```

```
# checking dtype success
```

```
df_movie_gross.dtypes
```

```
title           object
studio          object
domestic_gross  float64
foreign_gross   float64
year            int64
dtype: object
```

```
#Performing Visualization of df_movie_gross
```

```
income_foreign = df_movie_gross.foreign_gross
```

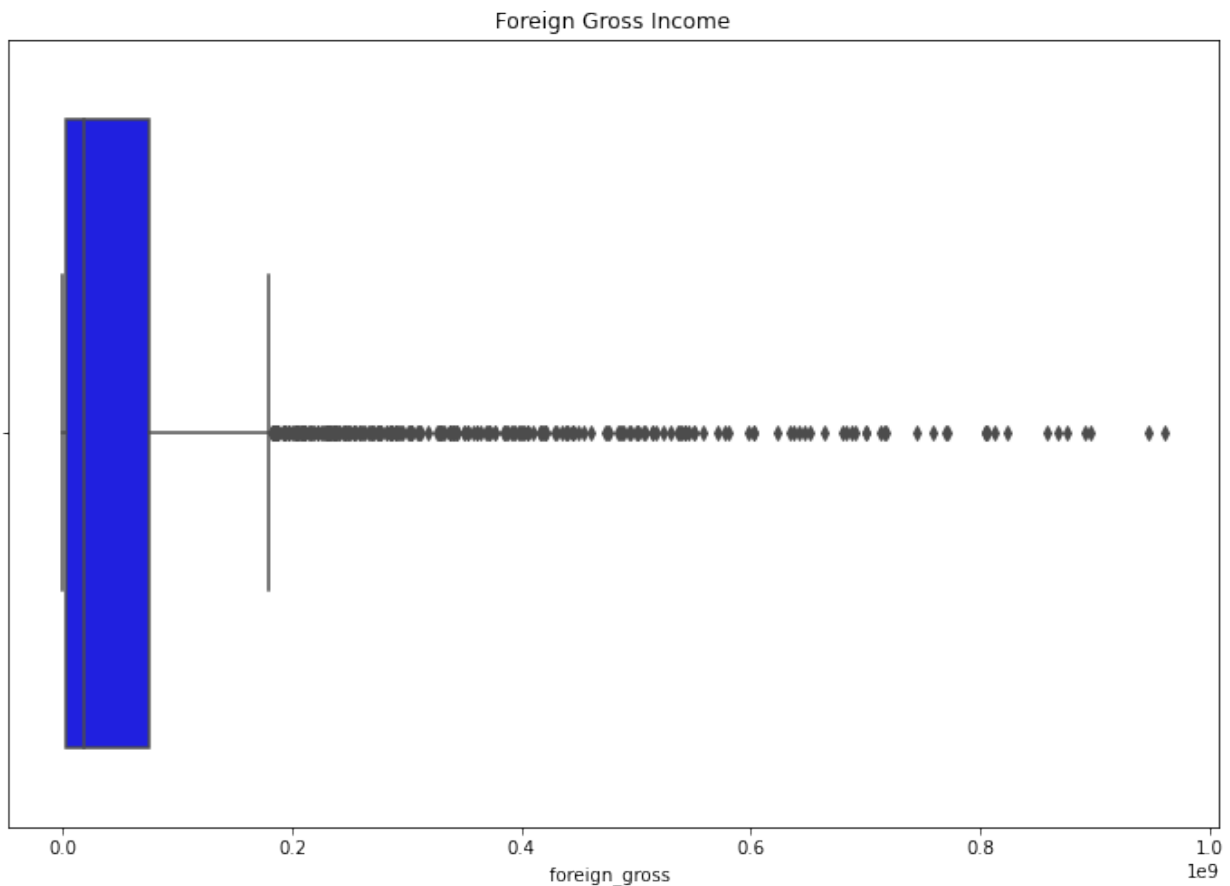
```
# find minimum and maximum values in runtime_minutes coumn
```

```
min_foreign = income_foreign.min()
max_foreign = income_foreign.max()
mean_foreign = income_foreign.mean()
print(f"minimum runtime: {min_foreign}")
print(f"Maximum runtime: {max_foreign}")
print(f"Mean runtime: {mean_foreign}")
# selecting the column for the boxplot
```

```
col_data = df_movie_gross.foreign_gross
# Creating boxplot using Seaborn Library
plt.figure(figsize= (12,8))
sns.boxplot(x= col_data, color= "blue")
plt.title('Foreign Gross Income')
plt.xlabel('foreign_gross ')
```

minimum runtime: 600.0
Maximum runtime: 960500000.0
Mean runtime: 74954901.2673389

Text(0.5, 0, 'foreign_gross ')



```
#Replacing the missing vlaue with median
df_movie_gross.foreign_gross.fillna(df_movie_gross.foreign_gross.isna(
).median, inplace=True)
```

```
#Checking whether the performance was successful
missing_values(df_movie_gross)
```

	index	Missing Values	Percentage
0	domestic_gross	28	0.008267
1	studio	5	0.001476

```

2         year                0    0.000000
3  foreign_gross              0    0.000000
4         title                0    0.000000

# Visualizing distribution of domestic_gross in our DataFrame
income_domestic = df_movie_gross.domestic_gross


# getting the minimum and maximum values in our runtime_minutes col
min_domestic = income_domestic.min()
max_domestic = income_domestic.max()
mean_domestic = income_domestic.mean()
print(f"minimum runtime: {min_domestic}")
print(f"Maximum runtime: {max_domestic}")
print(f"Mean runtime: {mean_domestic}")

# selecting the column for the boxplot
col_data = df_movie_gross.domestic_gross

# Creating boxplot & histogram using Seaborn Library
fig, ax = plt.subplots(ncols=2, nrows=1, figsize= (15,5))
sns.boxplot(x=col_data, ax=ax[0], color='yellow')
ax[0].set_title('Domestic Income')
ax[0].set_xlabel('domestic_gross')
df_movie_gross.domestic_gross.hist(ax=ax[1], color='blue')
ax[1].set_title('Histogram ')
ax[1].set_xlabel('domestic_gross ')
plt.tight_layout;

minimum runtime: 100.0
Maximum runtime: 936700000.0
Mean runtime: 28745845.06698422

```



```

#Repacng the missing values with Median
df_movie_gross.domestic_gross.fillna(df_movie_gross.domestic_gross.isn

```



```
a().median(), inplace=True)
```

```
#Verifying that the operation was successful
```

```
#function & checking for any instances of duplicates
```

```
print(f"Num of duplicates: {df_movie_gross.duplicated().sum()}")
```

```
print(missing_values(df_movie_gross))
```

```
Num of duplicates: 0
```

	index	Missing Values	Percentage
0	studio	5	0.001476
1	year	0	0.000000
2	foreign_gross	0	0.000000
3	domestic_gross	0	0.000000
4	title	0	0.000000

```
#Feature Engineering
```

```
#Merging different dataset in a single dataset
```

```
combined_data = df_movie_gross.merge(df_title_basics, left_on='title',  
right_on='original_title', how='left')
```

```
combined_data = combined_data.merge(df_title_ratings,  
left_on='tconst', right_on='tconst', how='left')
```

```
combined_data.head()
```

	title	studio	domestic_gross
0	Toy Story 3	BV	415000000.0
1	Alice in Wonderland (2010)	BV	334200000.0
2	Harry Potter and the Deathly Hallows Part 1	WB	296000000.0
3	Inception	WB	292600000.0
4	Shrek Forever After	P/DW	238700000.0

	foreign_gross	year	tconst	primary_title
0	652000000	2010	tt0435761	Toy Story 3
1	691300000	2010	NaN	NaN
2	664300000	2010	NaN	NaN
3	535700000	2010	tt1375666	Inception
4	513900000	2010	tt0892791	Shrek Forever After

	start_year	runtime_minutes	genres
averagerating	\		

```

0      2010.0      103.0  Adventure,Animation,Comedy
8.3
1      NaN      NaN      NaN
NaN
2      NaN      NaN      NaN
NaN
3      2010.0      148.0  Action,Adventure,Sci-Fi
8.8
4      2010.0      93.0  Adventure,Animation,Comedy
6.3

```

```

      numvotes
0  682218.0
1      NaN
2      NaN
3  1841066.0
4  167532.0

```

```

#Checking the missing value in the comnined data
missing_values(combined_data)

```

	index	Missing Values	Percentage
0	foreign_gross	1557	0.392686
1	numvotes	1522	0.383859
2	averagerating	1522	0.383859
3	genres	1225	0.308953
4	runtime_minutes	1225	0.308953
5	start_year	1225	0.308953
6	original_title	1225	0.308953
7	primary_title	1225	0.308953
8	tconst	1225	0.308953
9	domestic_gross	35	0.008827
10	studio	5	0.001261
11	year	0	0.000000
12	title	0	0.000000

```
combined_data.shape
```

```
(3965, 13)
```

```

#Accessing Columns of interest and reassinging to df_movies
df_movies = df_movies.loc[:, ['original_title', 'vote_average',
'vote_count', 'release_date']]

```

```

#Checking whether the needed data is successfully extracted
df_movies.head()

```

	original_title	vote_average
vote_count \		
0	Harry Potter and the Deathly Hallows: Part 1	7.7
	10788	

1	How to Train Your Dragon	7.7
7610		
2	Iron Man 2	6.8
12368		
3	Toy Story	7.9
10174		
4	Inception	8.3
22186		

	release_year
0	2010
1	2010
2	2010
3	1995
4	2010

```
# Convert 'release_date' column to date_time if it's not already in
datetime format
df_movies['release_date'] = pd.to_datetime(df_movies['release_date'])

# Extract the year from 'release_date' and create a new column
'release_year'
df_movies['release_year'] = df_movies['release_date'].dt.year

# Drop the 'release_date' column if you want to remove it
df_movies.drop('release_date', axis=1, inplace=True)

# # Display the first few rows of the modified DataFrame
print(df_movies.head())
```

	original_title	vote_average
vote_count \		
0	Harry Potter and the Deathly Hallows: Part 1	7.7
10788		
1	How to Train Your Dragon	7.7
7610		
2	Iron Man 2	6.8
12368		
3	Toy Story	7.9
10174		
4	Inception	8.3
22186		

	release_year
0	2010
1	2010
2	2010
3	1995
4	2010

#Executing an inner join to preserve only the movies appearing in both Tables

```
combined_data = combined_data.merge(df_movies, left_on= 'title',
right_on = 'original_title', how="inner")
combined_data.head()
```

	year \	title_x	studio	domestic_gross	foreign_gross
0	2010	Toy Story 3	BV	415000000.0	652000000
1	2010	Inception	WB	292600000.0	535700000
2	2010	Shrek Forever After	P/DW	238700000.0	513900000
3	2010	The Twilight Saga: Eclipse	Sum.	300500000.0	398000000
4	2010	Iron Man 2	Par.	312400000.0	311500000

	tconst	primary_title
0	tt0435761	Toy Story 3
1	tt1375666	Inception
2	tt0892791	Shrek Forever After
3	tt1325004	The Twilight Saga: Eclipse
4	tt1228705	Iron Man 2

	start_year	runtime_minutes	...	Unnamed: 0
0	2010.0	103.0	...	7
1	2010.0	148.0	...	4
2	2010.0	93.0	...	38
3	2010.0	124.0	...	15
4	2010.0	124.0	...	2

	id	original_language	original_title_y	popularity \
0	10193	en	Toy Story 3	24.445
1	27205	en	Inception	27.920
2	10192	en	Shrek Forever After	15.041
3	24021	en	The Twilight Saga: Eclipse	20.340

4	10138	en	Iron Man 2	28.515
---	-------	----	------------	--------

	release_date		title_y	vote_average	vote_count
0	2010-06-17		Toy Story 3	7.7	8340
1	2010-07-16		Inception	8.3	22186
2	2010-05-16	Shrek Forever After		6.1	3843
3	2010-06-23	The Twilight Saga: Eclipse		6.0	4909
4	2010-05-07		Iron Man 2	6.8	12368

[5 rows x 23 columns]

missing_values(combined_data)

	index	Missing Values	Percentage
0	numvotes	443	0.141669
1	averagerating	443	0.141669
2	runtime_minutes	241	0.077071
3	start_year	84	0.026863
4	genres	84	0.026863
5	original_title_x	84	0.026863
6	primary_title	84	0.026863
7	tconst	84	0.026863
8	year	0	0.000000
9	studio	0	0.000000
10	domestic_gross	0	0.000000
11	foreign_gross	0	0.000000
12	release_year	0	0.000000
13	vote_count	0	0.000000
14	original_title_y	0	0.000000
15	vote_average	0	0.000000
16	title	0	0.000000

dropping the missing values in start_year

combined_data.drop('start_year', inplace=True, axis=1)

#Filling missing values in averagerating column with vote_average column values

combined_data.averagerating.fillna(combined_data.vote_average, inplace=True)

#dropping vote_average columns

combined_data.drop('vote_average', axis=1, inplace=True)

#dropping vote_count columns and refilling the null values in num_votes with median

combined_data.drop('vote_count', inplace=True, axis=1)

#checking whether the vote_count colum has dropped successfully

combined_data.head()

year \	title_x	studio	domestic_gross	foreign_gross
0	Toy Story 3	BV	415000000.0	652000000
2010				
1	Inception	WB	292600000.0	535700000
2010				
2	Shrek Forever After	P/DW	238700000.0	513900000
2010				
3	The Twilight Saga: Eclipse	Sum.	300500000.0	398000000
2010				
4	Iron Man 2	Par.	312400000.0	311500000
2010				

tconst	primary_title	
original_title_x \		
0 tt0435761	Toy Story 3	Toy Story 3
1 tt1375666	Inception	Inception
2 tt0892791	Shrek Forever After	Shrek Forever After
3 tt1325004	The Twilight Saga: Eclipse	The Twilight Saga: Eclipse
4 tt1228705	Iron Man 2	Iron Man 2

runtime_minutes	genres	averagerating
numvotes \		
0 103.0	Adventure,Animation,Comedy	8.3
682218.0		
1 148.0	Action,Adventure,Sci-Fi	8.8
1841066.0		
2 93.0	Adventure,Animation,Comedy	6.3
167532.0		
3 124.0	Adventure,Drama,Fantasy	5.0
211733.0		
4 124.0	Action,Adventure,Sci-Fi	7.0
657690.0		

Unnamed: 0	genre_ids	id	original_language \
0 7	[16, 10751, 35]	10193	en
1 4	[28, 878, 12]	27205	en
2 38	[35, 12, 14, 16, 10751]	10192	en
3 15	[12, 14, 18, 10749]	24021	en
4 2	[12, 28, 878]	10138	en

	original_title_y	popularity	release_date	\
0	Toy Story 3	24.445	2010-06-17	
1	Inception	27.920	2010-07-16	
2	Shrek Forever After	15.041	2010-05-16	
3	The Twilight Saga: Eclipse	20.340	2010-06-23	
4	Iron Man 2	28.515	2010-05-07	

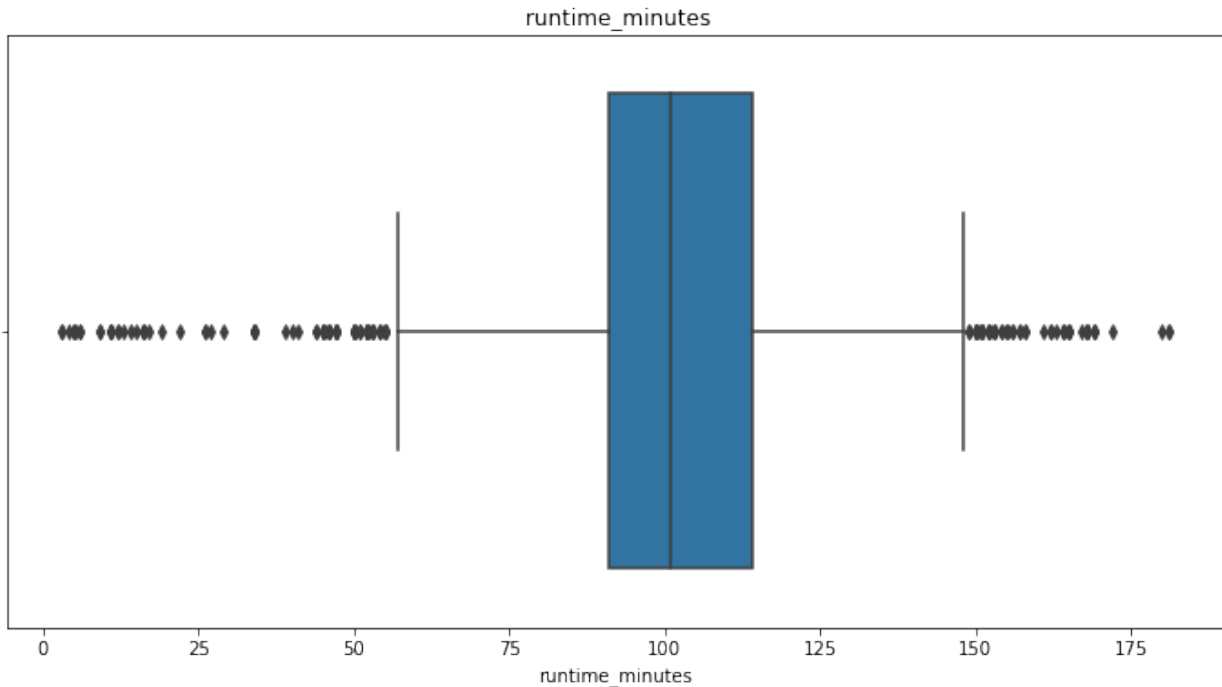
	title_y
0	Toy Story 3
1	Inception
2	Shrek Forever After
3	The Twilight Saga: Eclipse
4	Iron Man 2

missing_values(combined_data)

	index	Missing Values	Percentage
0	numvotes	1522	0.383859
1	averagerating	1522	0.383859
2	runtime_minutes	1364	0.344010
3	genres	1225	0.308953
4	original_title	1225	0.308953
5	primary_title	1225	0.308953
6	tconst	1225	0.308953
7	studio	5	0.001261
8	year	0	0.000000
9	foreign_gross	0	0.000000
10	domestic_gross	0	0.000000
11	title	0	0.000000

#Visualizing the distribution in combined dataframe for runtime_mintues

```
col_data = combined_data.runtime_minutes
plt.figure(figsize=(12,6))
sns.boxplot(x=col_data)
plt.title(' runtime_minutes');
```



```
#The data contains outliers
#Using imputation to replace the missing values
```

```
combined_data.runtime_minutes.fillna(combined_data.runtime_minutes.mean(),inplace=True)
```

```
#confirming if the operation was successful
missing_values(combined_data)
```

	index	Missing Values	Percentage
0	numvotes	444	0.141944
1	genres	85	0.027174
2	original_title_x	85	0.027174
3	primary_title	85	0.027174
4	tconst	85	0.027174
5	studio	1	0.000320
6	release_year	0	0.000000
7	original_title_y	0	0.000000
8	averagerating	0	0.000000
9	runtime_minutes	0	0.000000
10	year	0	0.000000
11	foreign_gross	0	0.000000
12	domestic_gross	0	0.000000
13	title	0	0.000000

```
#visualizing the distribution of numvotes
```

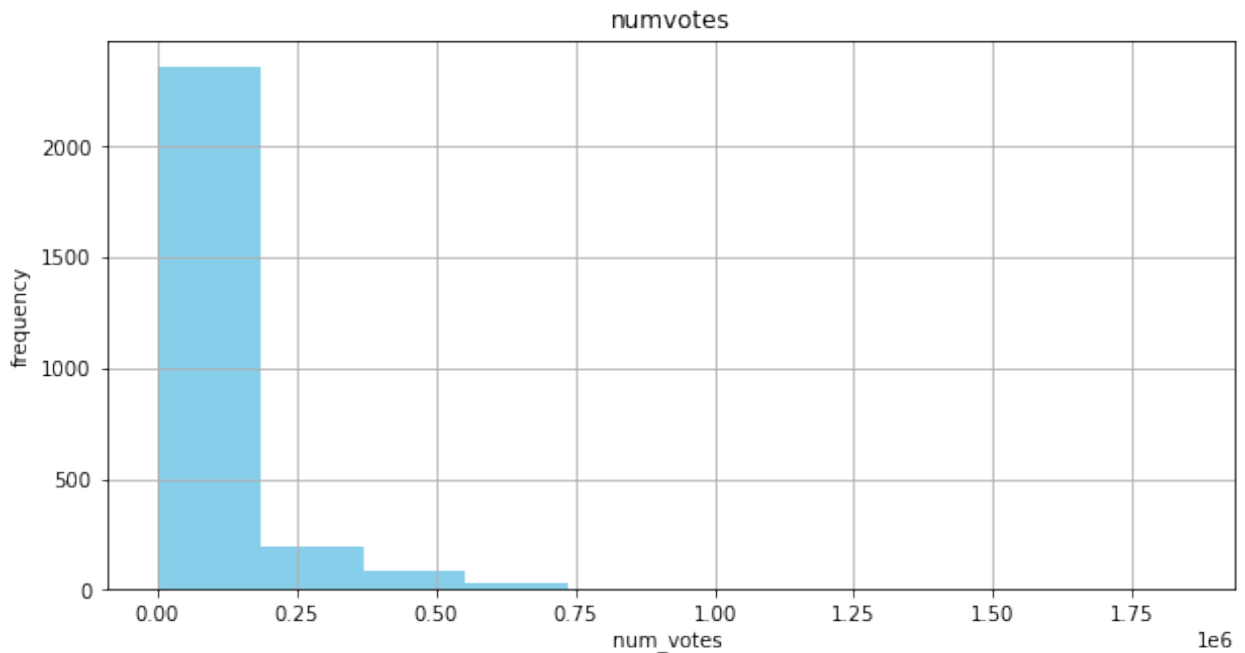
```
print(combined_data.numvotes.agg(['mean', 'std', 'min', 'max']))
combined_data.numvotes.hist(color='skyblue', figsize=(10, 5))
plt.title("numvotes")
```



```
plt.xlabel('num_votes ')
plt.ylabel('frequency')

mean    7.731767e+04
std     1.362478e+05
min     5.000000e+00
max     1.841066e+06
Name: numvotes, dtype: float64

Text(0, 0.5, 'frequency')
```



```
#Filling the skewed data using imputation median
combined_data.numvotes.fillna(combined_data.numvotes.median(),
inplace=True)
```

```
missing_values(combined_data)
```

	index	Missing Values	Percentage
0	genres	84	0.026863
1	original_title_x	84	0.026863
2	primary_title	84	0.026863
3	tconst	84	0.026863
4	release_year	0	0.000000
5	original_title_y	0	0.000000
6	numvotes	0	0.000000
7	averagerating	0	0.000000
8	runtime_minutes	0	0.000000
9	year	0	0.000000
10	foreign_gross	0	0.000000

11	domestic_gross	0	0.000000
12	studio	0	0.000000
13	title	0	0.000000

#Checking the most dominant genre using the process

#Calling the value_counts method to find the mode genre

```
top_genre = combined_data.genres.value_counts().head(1)
print(f'The most common genre in our data is: {top_genre}')
```

The most common genre in our data is: Drama 323
Name: genres, dtype: int64

#Replacing the missing values with mode imputation

```
combined_data.genres.fillna(combined_data.genres.mode().iloc[0],
inplace=True)
```

#confirming that we have no missing values in our dataset

```
missing_values(combined_data)
```

	index	Missing Values	Percentage
0	original_title_x	85	0.027174
1	primary_title	85	0.027174
2	tconst	85	0.027174
3	studio	1	0.000320
4	release_year	0	0.000000
5	original_title_y	0	0.000000
6	numvotes	0	0.000000
7	averagerating	0	0.000000
8	genres	0	0.000000
9	runtime_minutes	0	0.000000
10	year	0	0.000000
11	foreign_gross	0	0.000000
12	domestic_gross	0	0.000000
13	title	0	0.000000

#Confirming the dtypes of combined dataset

```
combined_data.dtypes
```

title	object
studio	object
domestic_gross	float64
foreign_gross	object
year	int64
tconst	object
primary_title	object
original_title_x	object
runtime_minutes	float64
genres	object
averagerating	float64
numvotes	float64

```
original_title_y    object
release_year        int64
dtype: object
```

```
#creating a new column that we will use to determine the financia
success of the company
```

```
combined_data['total_gross'] =
combined_data.domestic_gross.astype(str) +
combined_data.foreign_gross.astype(str)
```

```
#determining whether new column was created successfully
combined_data.head()
```

	title	studio	domestic_gross
0	Toy Story 3	BV	415000000
0	Toy Story 3	BV	415000000
0	Toy Story 3	BV	415000000
1	Alice in Wonderland (2010)	BV	334200000
2	Harry Potter and the Deathly Hallows Part 1	WB	296000000

	foreign_gross	year	tconst	primary_title	original_title	\
0	652000000	2010	tt0435761	Toy Story 3	Toy Story 3	
0	652000000	2010	tt0435761	Toy Story 3	Toy Story 3	
0	652000000	2010	tt0435761	Toy Story 3	Toy Story 3	
1	691300000	2010	NaN	NaN	NaN	
2	664300000	2010	NaN	NaN	NaN	

	runtime_minutes	genres	averagerating	numvotes
total_gross				
0	103.0	Adventure	8.3	682218.0
415000000652000000				
0	103.0	Animation	8.3	682218.0
415000000652000000				
0	103.0	Comedy	8.3	682218.0
415000000652000000				
1	NaN	NaN	NaN	NaN
334200000691300000				
2	NaN	NaN	NaN	NaN
296000000664300000				

```
combined_data['total_gross']
```

0	415000000652000000
0	415000000652000000
0	415000000652000000

```
1 334200000691300000
2 296000000664300000
```

```
...
3960 6200<bound method Series.median of 0 Fal...
3961 4800<bound method Series.median of 0 Fal...
3962 2500<bound method Series.median of 0 Fal...
3963 2400<bound method Series.median of 0 Fal...
3964 1700<bound method Series.median of 0 Fal...
```

```
Name: total_gross, Length: 7452, dtype: object
```

```
#Splitting movie genres and exploded to allow for genre-specific analyses
```

```
combined_data.genres = combined_data.genres.str.split(',')
combined_data.head()
```

	title	studio	domestic_gross
\			
0	Toy Story 3	BV	415000000.0
1	Alice in Wonderland (2010)	BV	334200000.0
2	Harry Potter and the Deathly Hallows Part 1	WB	296000000.0
3	Inception	WB	292600000.0
4	Shrek Forever After	P/DW	238700000.0

	foreign_gross	year	tconst	primary_title
original_title \				
0	652000000	2010	tt0435761	Toy Story 3
1	691300000	2010	NaN	NaN
2	664300000	2010	NaN	NaN
3	535700000	2010	tt1375666	Inception
4	513900000	2010	tt0892791	Shrek Forever After

	runtime_minutes	genres	averagerating
numvotes			
0	103.0	[Adventure, Animation, Comedy]	8.3
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	148.0	[Action, Adventure, Sci-Fi]	8.8

```
1841066.0
4          93.0  [Adventure, Animation, Comedy]          6.3
167532.0
```

#Exploding the genres

```
combined_data = combined_data.explode('genres')
```

```
combined_data.head()
```

	title	studio	domestic_gross	foreign_gross	year	
tconst \						
0	Toy Story 3	BV	415000000.0	6.52e+08	2010	tt0435761
0	Toy Story 3	BV	415000000.0	6.52e+08	2010	tt0435761
0	Toy Story 3	BV	415000000.0	6.52e+08	2010	tt0435761
1	Inception	WB	292600000.0	5.357e+08	2010	tt1375666
1	Inception	WB	292600000.0	5.357e+08	2010	tt1375666

	primary_title	original_title_x	runtime_minutes	genres
averagerating \				
0	Toy Story 3	Toy Story 3	103.0	Adventure
8.3				
0	Toy Story 3	Toy Story 3	103.0	Animation
8.3				
0	Toy Story 3	Toy Story 3	103.0	Comedy
8.3				
1	Inception	Inception	148.0	Action
8.8				
1	Inception	Inception	148.0	Adventure
8.8				

	numvotes	original_title_y	release_year
0	682218.0	Toy Story 3	2010
0	682218.0	Toy Story 3	2010
0	682218.0	Toy Story 3	2010
1	1841066.0	Inception	2010
1	1841066.0	Inception	2010

Plotting the domestic gross for each year


```
plt.figure(figsize=(10, 6))
plt.plot(combined_data['year'], combined_data['domestic_gross'],
marker='o', color='skyblue', linestyle='--')
```

Adding labels and title


```
plt.xlabel('Year')
plt.ylabel('Domestic Gross ($)')
plt.title('Domestic Gross for Each Year')
```

```
# Adding grid
plt.grid(True, linestyle='--', alpha=0.7)

# Show plot
plt.show()
```

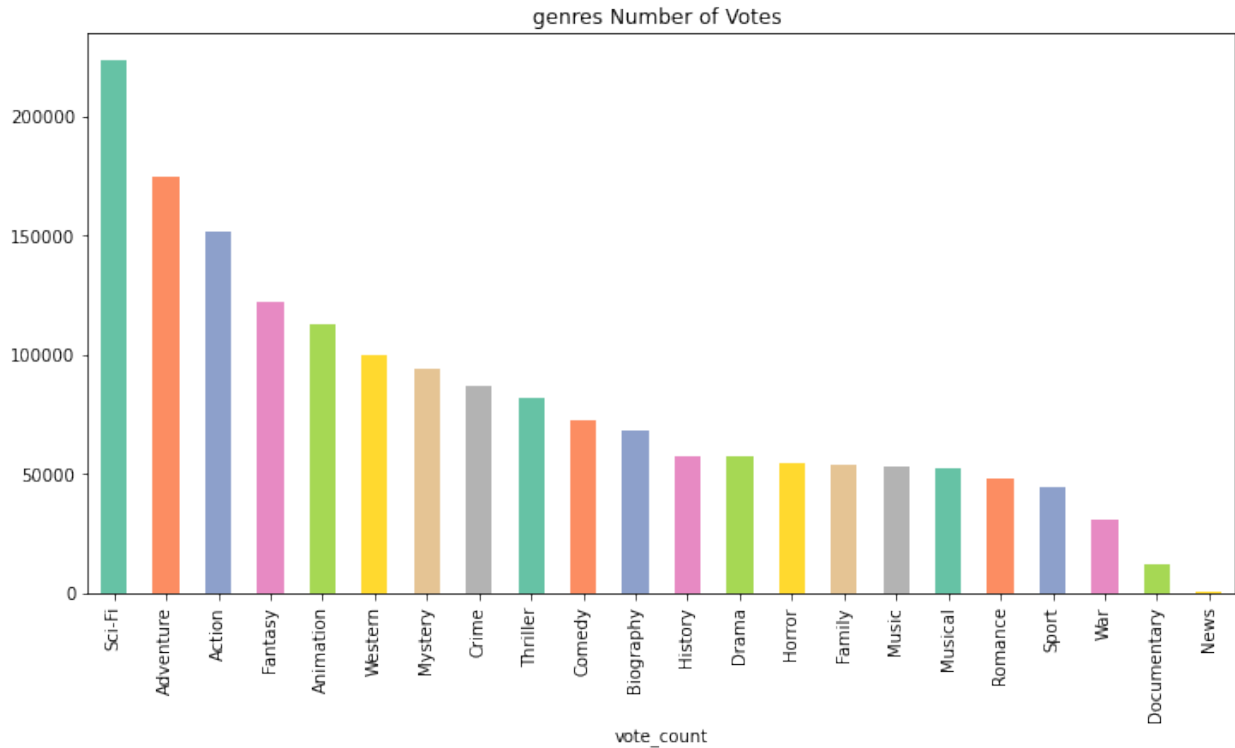


```
# Creating a heatmap to visualize the correlation matrix
plt.figure(figsize=(10, 8))
correlation_matrix = combined_data.corr()
plt.imshow(correlation_matrix, cmap= 'coolwarm',
interpolation='nearest')
plt.colorbar()
plt.xticks(range(len(correlation_matrix)), correlation_matrix.columns,
rotation = 45)
plt.yticks(range(len(correlation_matrix)), correlation_matrix.columns)
plt.title("Correlation Heatmap")
plt.show()
```



```
# Visualizing the relationship between genres and average ratings
plt.figure(figsize=(12, 6))
genre_avg_votes = combined_data.groupby('genres')['numvotes'].mean()
sorted_data = genre_avg_votes.sort_values(ascending=False)
colors = sns.color_palette("Set2", n_colors=len(sorted_data))
sorted_data.plot(kind='bar', color=colors)
plt.title("genres Number of Votes ")
plt.xlabel("vote_count")

Text(0.5, 0, 'vote_count')
```



```
# Visualizing the relationship between genre and average ratings
plt.figure(figsize=(12, 6))
genre_avg_ratings = combined_data.groupby('genres')
['averagerating'].mean()
sorted_data = genre_avg_ratings.sort_values(ascending=False)
sorted_data.plot(kind='bar', color='green')
plt.title("Average Ratings ")
plt.xticks(rotation=90)
plt.xlabel("Genre")
plt.ylabel("Average Rating");
```